

# Synthetic Data Generation using Benerator Tool

Vanessa Ayala-Rivera<sup>\*1</sup>, Patrick McDonagh<sup>†2</sup>, Thomas Cerqueus<sup>‡1</sup>  
and Liam Murphy<sup>§1</sup>

<sup>1</sup>Lero@UCD, Performance Engineering Laboratory, School of  
Computer Science and Informatics, University College Dublin

<sup>2</sup>Lero@DCU, Performance Engineering Laboratory, School of  
Electronic Engineering, Dublin City University

October 29, 2013

University College Dublin, Technical Report UCD-CSI-2013-03

## Abstract

Datasets of different characteristics are needed by the research community for experimental purposes. However, real data may be difficult to obtain due to privacy concerns. Moreover, real data may not meet specific characteristics which are needed to verify new approaches under certain conditions. Given these limitations, the use of synthetic data is a viable alternative to complement the real data. In this report, we describe the process followed to generate synthetic data using Benerator, a publicly available tool. The results show that the synthetic data preserves a high level of accuracy compared to the original data. The generated datasets correspond to microdata containing records with social, economic and demographic data which mimics the distribution of aggregated statistics from the 2011 Irish Census data.

## 1 Introduction

The creation of synthetic data is an approach widely used by the research community in a variety of domains: privacy protection [8, 6], healthcare [5], pattern recognition [9], data mining [7], etc. Such data is often generated to meet specific characteristics that are not found in the real data. By generating synthetic datasets, researchers can have more flexibility on the manipulation of the data and are able to test a wider set of conditions and scenarios in their applications.

---

<sup>\*</sup>vanessa.ayala-rivera@ucdconnect.ie

<sup>†</sup>patrick.mcdonagh@dcu.ie

<sup>‡</sup>thomas.cerqueus@ucd.ie

<sup>§</sup>liam.murphy@ucd.ie

Moreover, synthetic data is also used as a substitute for real data, as it is often difficult to obtain due to privacy concerns (i.e., to protect the privacy of the individuals represented in the real data).

It is often the case that researchers adopt the practice of generating synthetic data to test their proposed algorithms against more heterogeneous datasets. In particular, some of the reasons why we believe the generation of synthetic data is a useful approach for the research community are the following:

- Generating synthetic data allows to control the data distributions used for testing. One can study the behavior of the algorithms under different conditions: identify scenarios where the data distribution favors the performance of the algorithm or the scenarios where the algorithm performs the worst.
- Synthetic data can help to allow a fair performance comparison among the algorithms. For example, for evaluating the scalability of the algorithms. When the same dataset is used for testing, but only scaling up its size while preserving the same data distribution, the measures obtained in terms of efficiency (e.g., memory consumption, running time) give a more precise idea about the causes for the increase of the computational resources. Especially for algorithms that could be affected by some aspects of the datasets like the cardinality of the attributes.
- Generating synthetic data allows to create records which have the finest level of granularity in each attribute. In contrary, publicly available real datasets have often undergone anonymization procedures due to privacy constraints. Therefore, the values for some of the attributes are already grouped in less specific values. For example, when some values are sparse in the dataset, they are all placed together in a group to protect the privacy of the individuals that fall in that minority population.

In order to benefit from this approach, it is important to have practical tools that can be customized and easily extended according to different needs.

In this work we present a practical approach to generate synthetic census-based data. We describe the methodology followed to generate this type of datasets and how to simplify its creation using Benerator [4], a publicly available tool.

## 2 Synthetic Data Generation

In this section we describe the process followed to generate microdata [10], where the records correspond to the information about an individual. The process can be applied to microdata of different domains. However, in our work, we have applied it to create personal data records. Each record contains information about social, economic and demographic attributes of a person. This data is based on attributes collected from the Irish Census 2011<sup>1</sup>.

---

<sup>1</sup><http://www.cso.ie/en/databases/>

Given that access to microdata is commonly restricted, the census publishes only the aggregated statistics for some of the attributes. Often, these attributes are correlated with each other and the statistics for these relationships are disseminated. One example of the published relationships are: marital status by age group; and highest level of education completed by socio-economic group. In our analysis of census data, we capture the frequency distributions of multiple attributes and use them as probability weights in the data generation process. By using this approach, we preserve the density of the population corresponding to the selected demographic attributes.

We captured the statistics from people corresponding to the adult population only (i.e., ages between 17 and 84 years). This selection results in a population count of 3,550,246 people. This is the total number of records we considered as the original Irish Census data in our work to verify the accuracy in which the synthetic data is generated.

## 2.1 Overview of the Generation Process

There are different frameworks to generate synthetic data [1, 2, 3]. In our case we used the Java open-source tool Benerator [4], as it offers high capabilities of extensibility and customization. We implemented the logic to use the census data as the domain values and the aggregated statistics as weights for the generation process.

Figure 1 shows the steps and components involved in the generation of the synthetic data. Firstly, we configure a descriptor file, which is stored in XML format. In this file we indicate the number of records to be generated, the type of entity to be created (e.g., a census-based person) and its attributes (e.g., *age*, *gender*, *marital status*); and the type of output for the dataset (e.g., plain text file, databases). In a separate file, which is stored in comma-separated value format (CSV), we also configure the universe of values that each attribute can take and the distribution they will follow in the dataset. Whereas the records can be randomly generated, in our case, we wanted to mimic the distribution from the aggregated statistics from the Irish Census.

Once the metadata has been set up, the generator object for the entity (i.e., *entity generator*) will start the creation of the synthetic records. The *Entity generator* acts as the controller to run the *sub generators* corresponding to each of the demographic attributes: Once the entity object is created, it will be made available to the other *sub generators* so they can use it in their generation process. The *entity generator* has also the logic to respect a set of constraints between some of the attributes (one of the main characteristics of realistic datasets). This logic ensures to assign realistic values for the attributes. For example, when generating *marital status*, it is important to consider that most of the population between 15 and 19 years are single thus the *generator* should not assign a widowed status, which is less likely.

The outcome of the synthetic generation process is a set of entities (e.g., census-based person records) which are stored in CSV files. We chose to use CSV format because it is widely-known format, easy to manipulate (i.e., ex-

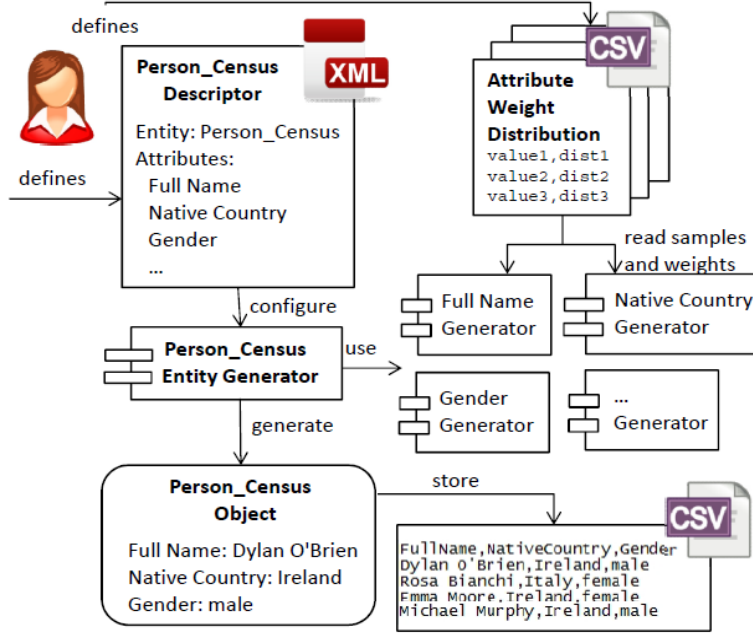


Figure 1: Overview of the synthetic data generation process.

port/import) and portable among different types of applications (e.g., supported by most of frameworks and tools).

## 2.2 Implementation Details

We customized Benerator to generate personal records according to the Irish Census data. Figure 2 shows the classes involved in our customization. In order to respect the multi-attribute constraints, we followed a specific order in the generation of the data that is part of the *personcensus* entity. This logic is implemented in the *PersonCensusGenerator* class. The first attributes to be generated are the independent ones: *gender*, *nationality* and *age*. Once the values for those attributes have been generated, the dependent attributes can use them to generate their own data. For example, *full name* is dependent on *gender* and *nationality* to be customized accordingly. *Industrial group* and *field of study* are dependent on *age* and *gender*. The rest of the attributes (*education*, *marital status*, etc.) are generated using *age* groups as base.

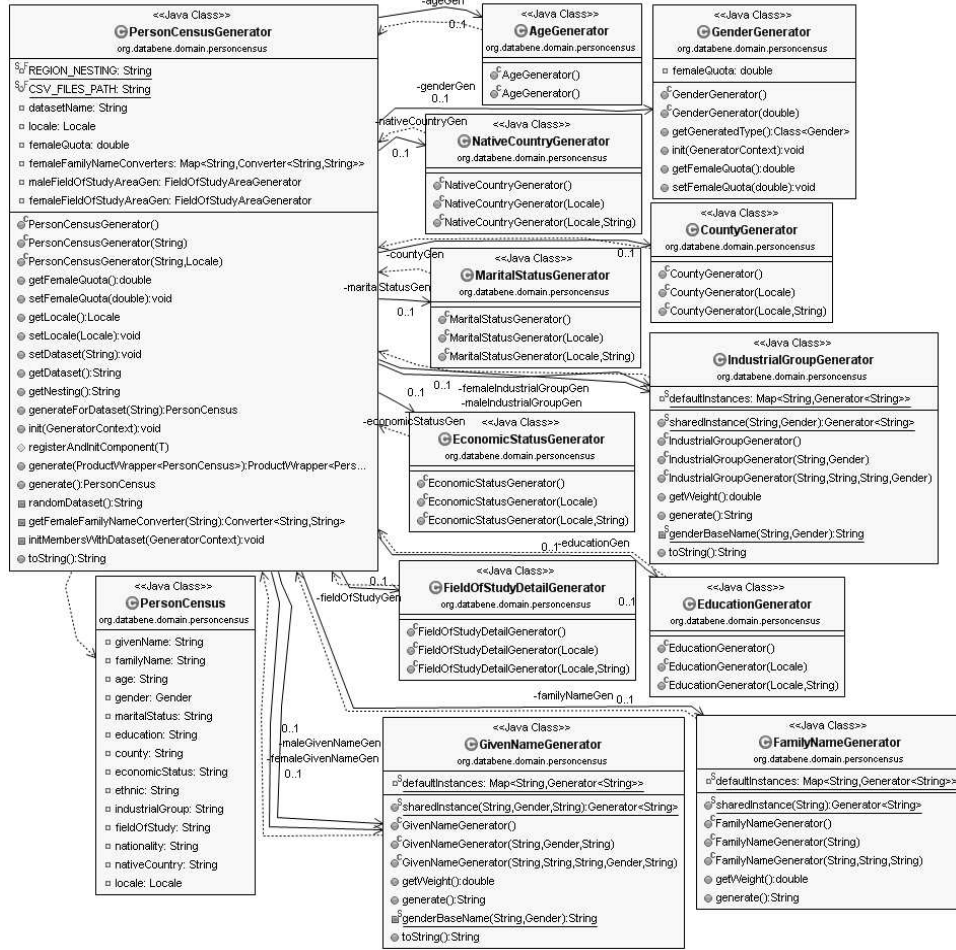


Figure 2: Class diagram for the generation of person census records.

To better illustrate the synthetic generation process, we provide an example of the configuration files we defined in our work.

**Descriptor file:** We configured the location as *IE* to indicate Benerator to use the set of files regionalized for *Ireland*. The type of entity to generate is setup as *PersonCensus* which is the name for the generator class. The number of entities to be created is set to 30,000. The generated entity is kept in a variable with name *person*. All the attributes that are part of this entity are listed with their corresponding sub generator. Finally, the output is directed to a CSV file. An example of a descriptor file is shown in *PersonCensusDescriptor.xml*.

---

```
// PersonCensusDescriptor.xml
<?xml version="1.0" encoding="iso-8859-1"?>
<setup defaultDataset="IE">
  <import domains = "personcensus" />
  <generate type="PersonCensus" count="30000" >
    <variable name="person" generator="PersonCensusGenerator"
      dataset="IE" locale="IE"/>
    <attribute name="gender" script="person.gender" />
    <attribute name="age" script="person.age" />
    <attribute name="maritalStatus" script="person.maritalStatus" />
    <attribute name="economicStatus" script="person.economicStatus"/>
    <attribute name="NativeCountry" script="person.nativeCountry" />
    <attribute name="FullName" script="person.givenName + ' ' +
      person.familyName" />
    <consumer class="org.databene.platform.csv.CSVEntityExporter">
      <property name="uri" value="./output/irishcensus30m.csv"/>
      <property name="columns" value="FullName, NativeCountry,
        Gender, Age, MaritalStatus, EconomicStatus"/>
    </consumer>
  </generate>
</setup>
```

---

**Configuration CSV files:** We defined the universe of values for each attribute and the weights to be used by the distribution function in Benerator. For the independent attributes, like *nationality*, a single file is configured (as shown in file *Nationality.csv*). In contrary, for dependent attributes, like *marital status* which depends on *age*, one file is configured for each *age group* (as shown in files *MaritalStatusQty15-19.csv* and *MaritalStatusQty80-84.csv*).

---

```
//Nationality.csv
Irish,969087
Austrian,554
Belgian,932
Bulgarian,865
Cypriot,53
Czech,1921
...
Australian,2190
New Zealander,914

//MaritalStatusQty15-19.csv
Single,282106
Married (first marriage),836
Re-married (following widowhood),4
Re-married (following dissolution of previous marriage),3
Separated (including deserted),55
Divorced,8
Widowed,7
```

```
//MaritalStatusQty80-84.csv
Single,11898
Married (first marriage),25133
Re-married (following widowhood),565
Re-married (following dissolution of previous marriage),209
Separated (including deserted),699
Divorced,308
Widowed,31301
```

---

**Generated Dataset (CSV file):** In Figure 3, we show an example of a dataset generated based on the Irish Census.

FullName	NativeCountry	Gender	Age	MaritalStatus
Fionn Murphy	Ireland	male	73	Married (first marriage)
Ying Dong	China	female	48	Married (first marriage)
Paul O'Kelly	Ireland	male	76	Widowed
Lucy O'Neill	Ireland	female	17	Single
Megan O'Reilly	Ireland	female	83	Married (first marriage)
Maja McCarthy	Ireland	female	68	Married (first marriage)
Grace Gallagher	Ireland	female	53	Married (first marriage)
Mia O'Neill	Ireland	female	61	Single
Michael Ryan	Ireland	male	56	Married (first marriage)
Susanne Walter	Germany	female	36	Married (first marriage)
Alfie O'Reilly	Ireland	male	40	Married (first marriage)
Ruairi Smith	Ireland	male	29	Married (first marriage)
Daniel O'Byrne	Ireland	male	29	Married (first marriage)

Figure 3: Example of generated data for person census domain.

### 3 Results

In this section we describe the generated datasets: detailing the attributes which are part of the datasets and the domain values for each attribute. Moreover, we show the comparison between the data distributions from the Irish Census data and the distribution from a dataset that was synthetically generated based on the aggregate statistics of the census data. The original census dataset was composed of 3,550,246 records and we scaled down the dataset population to different sizes. The results shown in this report are those obtained for a dataset with 100k records.

### 3.1 Comparison of Real and Synthetic Datasets

#### 3.1.1 Independent Attributes

Figure 4 shows the comparison between the distribution from the Irish Census data (a) and the synthetic dataset (b) for the *age* attribute. It can be seen that the data distribution between the synthetic and the original data are similar, showing that the generation process preserves a good level of accuracy from the original distribution.

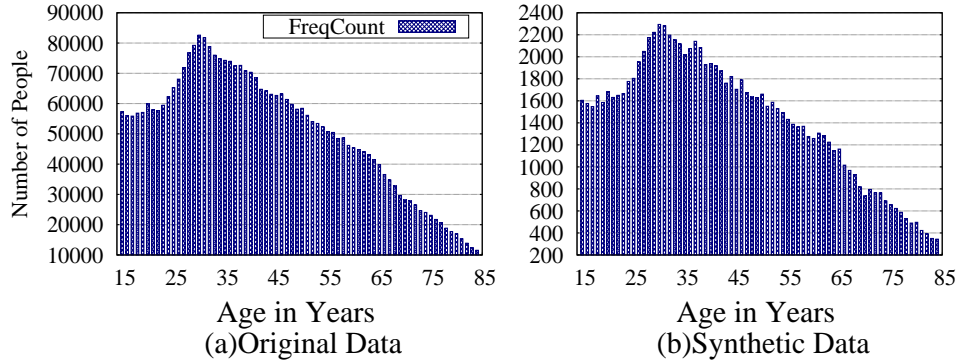


Figure 4: Comparison of data distribution between original and synthetic dataset for the *age* attribute.

#### 3.1.2 Dependent Attributes

Figure 5 shows the comparison between the distribution from the Irish Census data (a) and the synthetic dataset (b) for the *marital status* attribute, which is constrained by *age group*. It can be seen that the generated synthetic data preserves a high level of accuracy compared to the original distribution among the *age* groups.

## 4 Synthetic Datasets Description

The generated datasets are formed by the following attributes:

1. *Full name*: These are randomly generated names that result from a concatenation of a list of given names and family names. Full names are generated according to the nationality and age attributes to match the most commonly used names in those countries.
2. *Age*: Values between 17 and 84 in order to only consider the adult population.
3. *Gender*: female, male.



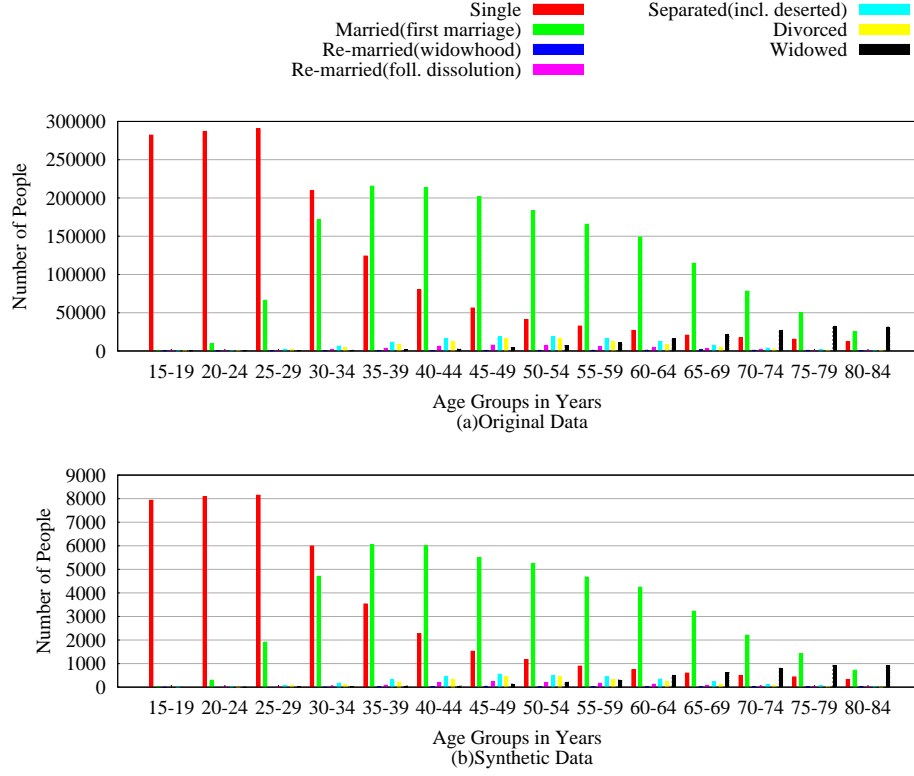


Figure 5: Comparison of data distribution between original and synthetic dataset for constraint attribute *marital status per age groups*.

4. *County*: The list of all Irish counties which are: Carlow, Dublin City, Dn Laoghaire-Rathdown, Fingal, South Dublin, Kildare, Kilkenny, Laois, Longford, Louth, Meath, Offaly, Westmeath, Wexford, Wicklow, Clare, Cork City, Cork County, Kerry, Limerick City, Limerick County, North Tipperary, South Tipperary, Waterford City, Waterford County, Galway City, Galway County, Leitrim, Mayo, Roscommon, Sligo, Cavan, Donegal, Monaghan.
5. *Marital Status*: Single, Married (first marriage), Re-married (following widowhood), Re-married (following dissolution of previous marriage), Separated (including deserted), Divorced, Widowed.
6. *Native Country*: Ireland, Austria, Belgium, Bulgaria, Cyprus, Czech Republic, Denmark, Estonia, Finland, France, Germany, Greece, Hungary, Italy, Latvia, Lithuania, Luxembourg, Malta, Netherlands, Poland, Portugal, Romania, Slovakia, Slovenia, Spain, Sweden, Russian Federation, Ukraine, Niger, South Africa, Mauritius, India, Philippines, China, Pak-

istan, Malaysia, United States of America, Brazil, Canada, Australia, New Zealand.

7. *Economic Status*: Employer or own account worker, Employee, Assisting relative, Unemployed looking for first regular job, Unemployed having lost or given up previous job, Student or pupil, Looking after home/family, Retired, Unable to work due to permanent sickness or disability, Other economic status.
8. *Industrial Group*: Agriculture, forestry and fishing (A), Mining and quarrying (B), Manufacturing (C), Electricity, gas, steam and air conditioning supply (D), Water supply; sewerage, waste management and remediation activities (E), Construction (F), Wholesale and retail trade; repair of motor vehicles and motorcycles (G), Transportation and storage (H), Accommodation and food service activities (I), Information and communication (J), Financial and insurance activities (K), Real estate activities (L), Professional, scientific and technical activities (M), Administrative and support service activities (N), Public administration and defence; compulsory social security (O), Education (P), Human health and social work activities (Q), Arts, entertainment and recreation (R), Other service activities (S), Activities of households as employers producing activities of households for own use (T), Activities of extraterritorial organisations and bodies (U).
9. *Education*: No formal education, Primary, Lower secondary, Upper secondary, Technical/vocational, Advanced certificate/completed apprenticeship, Higher certificate, Ordinary bachelor degree/professional qualification or both, Honours bachelor degree/professional qualification or both, Postgraduate diploma or degree, Doctorate (Ph.D).
10. *Field of Study*: Education and teacher training, Music and performing arts, Audio-visual techniques and media production, Design, Other arts, Foreign languages, Mother tongue, History and archaeology, Other humanities, Psychology, Economics, Business and administration (broad programmes), Marketing and advertising, Accounting and taxation, Management and administration, Secretarial and office work, Law, Other social sciences, business and law subjects, Biology and biochemistry, Physical sciences (physics, chemistry, earth science), Computer science, Computer use, Other science, mathematics and computing, Engineering and engineering trades (broad programmes), Mechanics and metalwork, Electricity and energy, Motor vehicles, ships and aircraft, Architecture and town planning, Building and civil engineering, Other engineering, manufacturing and construction, Crop and livestock production, Other agriculture and veterinary, Medicine, Nursing and caring, Child care and youth services, Social work and counselling, Other health and welfare, Hotel, restaurant and catering, Hair and beauty services, Other personal services, Air

transportation, Ground transportation, Sea transportation, Other transportation services, Public security services, Industrial security services, Other security services, Other subjects.

## 5 Conclusions

In this report we described the approach we used to generate synthetic micro-data. We described how we extended the public available tool, called Benerator, to add a new domain for data generation using census-based personal records. We showed how capturing the frequency distributions from a census dataset and using them as probability weights in the generation process, we were able to mimic the data distribution for that census. We also presented the comparison of the distributions between the original and the synthetic data, showing that synthetic dataset preserves a high level of accuracy compared to the original distribution. Finally, we provided a detailed description of the generated datasets about the domain values used for each attribute.

## Acknowledgements

This work was supported, in part, by Science Foundation Ireland grant 10/CE/I1855 to Lero - the Irish Software Engineering Research Centre ([www.lero.ie](http://www.lero.ie)). This work has also received support from Science Foundation Ireland (SFI) via grant 08/SRC/I1403 FAME SRC (Federated, Autonomic Management of End-to-End Communications Services - Scientific Research Cluster).

## References

- [1] DBMonster. <http://dbmonster.sourceforge.net>.
- [2] generatedata.com <http://www.generatedata.com>.
- [3] Jailer. <http://jailer.sourceforge.net/home.htm>.
- [4] Volker Bergmann. Data Benerator Tool. <http://databene.org/databene-benerator>.
- [5] Philip Cooley, Bruce Y Lee, Shawn Brown, James Cajka, Bernadette Chasteen, Laxminarayana Ganapathi, James H Stark, William D Wheaton, Diane K Wagener, and Donald S Burke. Protecting health care workers: a pandemic simulation based on Allegheny County. *Influenza and other respiratory viruses*, 4(2):61–72, March 2010.
- [6] Gabriel Ghinita, Yufei Tao, and Panos Kalnis. On the Anonymization of Sparse High-Dimensional Data. In *ICDE*, pages 715–724. Ieee, April 2008.

- [7] Guimei Liu, H Zhang, and L Wong. Controlling false positives in association rule mining. *Proceedings of the VLDB Endowment*, pages 145–156, 2011.
- [8] Ashwin Machanavajjhala, Daniel Kifer, John Abowd, Johannes Gehrke, and Lars Vilhuber. Privacy: Theory meets Practice on the Map. *ICDE*, pages 277–286, April 2008.
- [9] T Varga and Horst Bunke. Generation of synthetic training data for an HMM-based handwriting recognition system. *Document Analysis and Recognition*, 2003.
- [10] LCRJ Willenborg and T De Waal. Elements of statistical disclosure control. *Lecture Notes in Statistics*, 155, 2001.